

Manipulating Model Explanations: Why you shouldn't trust me

Verena Heusser
verena.heusser@student.kit.edu
Karlsruhe Institute of Technology (KIT)
Karlsruhe, Germany

ABSTRACT

This paper reviews state-of-the-art approaches to model explanations with a focus on those techniques that try to fool these methods.

We don't only want the model to be good. We want it to be safe and interpretable. As machine learning models enter critical areas in human lives such as the criminal justice system, medicine or financial systems, the inability for humans to understand these models is dangerous and problematic. Advances in the rising research area of explainable AI seems to be a remedy. However, there is not yet a consensus about the validity and robustness of explanations methods themselves. The main suggestion of this paper is to be cautious about results of explanation methods. Explanations can be fooled just as the underlying machine learning models. So, in the end the question must be posed whether inexplainable models should be used at all if we need other models to explain these models but are not valid themselves ..

KEYWORDS

Interpretability, Neural networks, adversarial training

ACM Reference Format:

Verena Heusser. 2020. Manipulating Model Explanations: Why you shouldn't trust me. In *Explainability '20: ACM Symposium on Neural Model Explanations*, December 16–20, 2020, Karlsruhe, DE. ACM, New York, NY, USA, 6 pages. <https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

1 INTRODUCTION

In recent years, deep learning models have demonstrated superior performance on a variety of tasks [3, 11, 13]. While performance is still increasing and more tasks are being handled, their performance comes at the cost of complexity: models often use millions to billions of parameters to achieve universal function approximation. This complexity means that they remain black boxes that cannot be interpreted even by experts.

Such a black box is able to predict well for unseen yet similar data, answering the question of *what* is the most likely label for an input sample. However, most models will provide no answer to *why* the model chose this label for the instance and which features of the instance were crucial for making this prediction.




Figure 1: Prediction pipeline using a machine learning model, depicted as black box. Typically, evaluation metrics require the prediction y and the ground truth label y^* allowing for the assessment of the model's accuracy. Additionally answering the question *why*, i.e. making the model interpretable for a human, requires additional methods.

For instance, withing object recognition one would likely assume that the presence (or absence) of an object in the image causes a model to decide for a specific object category, closely akin to how humans base their decision process. As seen in Figure 1, the presence of the cat in the input image x should be causing the model to predict the label 'cat'.

Automated algorithms are already in use in critical areas, such as medicine TODOcite, the criminal justice system [4] and the financial sector TODOcite or the piloting of self-driving cars. Thus, as machine learning models are moving out of the lab into the real-world, the inability of humans to understand these models seems even more problematic.

Thus, automated interpretation methods are required to make sense of the reasoning process and robustness of such deep learning based models and to ensure that a model makes decisions without unfair or hidden biases. The research field approaching the explanation or validation of machine learning models is called explainable artificial intelligence (XAI).

Not knowing about the biases of a network the vastly advancing technology of machine learning to be used in high-stakes and safety critical applications and prevent real-life deployment of such systems. Furthermore, the rise in machine learning model deployments also caused the development of adversarial attacks. These attacks attempt to fool a machine learning model by providing deceptive input. Fooling refers to the resulting malfunction of the model.

In particular, it is unclear how the variety of proposed model interpretation methods are related and what common concepts can be used to evaluate and compare them. Many works are dedicated to establish a formal definition of what it means to be interpretable, and how to select, evaluate and compare methods for generating explanations of machine learning models. [9, 10].

Not knowing about attacks and data arranged to exploit specific vulnerabilities has contributed to a relatively new research field

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

Explainability '20, December 16–20, 2020, Karlsruhe, DE

© 2020 Association for Computing Machinery.

ACM ISBN 978-x-xxxx-xxxx-x/YY/MM...\$15.00

<https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

of XAI comprising topics of (1) *model explanations*, (2) *adversarial attacks*, or manipulation methods and (3) the field of *adversarial manipulations of model explanations*. All of this is also known by the name of robust machine learning or even explainable artificial intelligence, as all subfields have the common goal to make models more robust and safe for deployment. (1) refers to the development of techniques that can be used to understand and explain the decision making process of a machine learning model or even the development of models that are inherently interpretable. (2) is the field of detecting vulnerabilities in models that cause models to be deceived by altered input. (3) is the main topic of this paper, i.e. how to fool explanation models in order to detect vulnerabilities and malfunctions in explanation methods.

Detecting such vulnerabilities in models is most crucial

Fragility limits how much we can trust and learn from the interpretations

Most research on explainability focuses on the application of computer vision tasks.

Most works in the field of XAI focus on image classification tasks, mostly because visualizations of a neural networks prediction can be easily verified by a human. The general purpose of image classification is to detect what objects are in an image. If a model works can be checked rather easily (if an image contains a cat, the prediction of a neural network should be cat and not some other object category). However, how it works (*interpretability*), i.e. based on which features in the image the decision is made or which parameters in the model influence the prediction most, is an entirely different matter (*explainability*).

More importantly, while a big motivation for the development of robust and explainable systems is to overcome biases in models, datasets with direct implication of biases are seldomly used and by far not treated as benchmarking scenarios for explainability analyses.

The overview presented in this article examines the existing literature and contributions in the field of XAI focusing on methods to manipulate explanation methods. The critical literature analysis might serve as a motivation and step towards the biggest problems in XAI: How to make sure that interpretations of models are truly valid. This paper is structured as follows...

2 INTERPRETATION METHODS

There exists a variety of definitions in the vastly expanding research field of XAI, and the concept of *interpretability* still has no formal commonly used technical meaning [9]. To build on the common ground of existing research, this paper follows the terminology of [2]. The authors make a distinction between the related but different concepts of *interpretability* and *explainability*.

Interpretability refers to the extent to which cause and effect can be observed in a model, which can be translated to uncovering *why* a model made prediction y to an input x . Consider the example of image classification from Figure 1. Here, interpretability would mean that if a cat is present in the image (the cause), the model classifies it to the category 'cat' (the effect). Now imagine we find that the model takes the green meadow in the image as evidence to predict 'cat', and not the cat itself, there is a lack of interpretability, as other features are learned as important than the features that

are truly important. This toy example is in fact a common problem in image classification: the over-reliance of models on image background, rather than on objects in the foreground, has been observed for the task of image classification [18].

Explainability on the other hand spans methods to uncover *how* a model makes predictions, meaning to observe the inner workings of a model and to literally explain what is happening.

While both concepts seem to be important for the general objective of explainable artificial intelligence, this paper focuses on interpretability. High interpretability is desired as it can help to uncover biases in the model. Suppose a machine learning model is to be deployed for the task of income prediction based on features such as age, race, gender, education and hours of work per week. The performance of the system would mainly be evaluated in terms of the predictive accuracy and the fairness of the system. The former can be evaluated with metrics, such as accuracy on a held-out test set. For the latter, interpretability methods might be applied to observe which input features are used by the model to predict the income. If the model uses sensitive features, such as sex and race as important features, it is systematically biased and thus unfair.

Explanation methods aim at making complex and inherently uninterpretable black box models interpretable by creating human readable visualizations. A frequently used type of explanation methods are feature attributions mapping each input feature to a numeric score. This score should quantify the importance of the feature relative to the model output. The resulting attribution map is then visualized as a heatmap projected onto the input sample to interpret the input attributes regarding which ones are the most helpful for forming the final prediction.

Definition 1: Interpretation Method

We consider a neural network $N : \mathbb{R}^d \rightarrow \mathbb{R}^k$. For the task of image classification, N classifies an input image $\mathbf{x} \in \mathbb{R}$ in k categories where the prediction $f_N(\mathbf{x}) = y \in \{1, \dots, K\}$ is given by $y = \operatorname{argmax}_i f_N(\mathbf{x})_i$.

Given the neural network N , its input vector $\mathbf{x} \in \mathbb{R}^d$ and the neural networks prediction for input \mathbf{x} , $f_N(\mathbf{x}) = y$, an interpretation method \mathcal{I} determines why label y has been chosen by N for input \mathbf{x} . The interpretation is given by an output vector $\mathbf{x} \in \mathbb{R}^d$ where each entry h_i is mostly a numeric value describing the relevance of an input dimension x_i of \mathbf{x} for $f_N(\mathbf{x})$. As \mathbf{x} has the same dimensions as the input \mathbf{x} it can be mapped to the input, overlaying \mathbf{x} as a heatmap, where the color value represents the importance of feature x_i towards the prediction $f_N(\mathbf{x})$.

An example is given in TODO. Higher values, implying a stronger relative importance for making the prediction $f_N(\mathbf{x})$, is depicted in TODD color.

While all explanation methods try to obtain importance measures for the network prediction, they differ with respect to how these measures are obtained. [17] propose two major categories for explanation strategies, namely *black-box* explanation methods and *white-box* explanation methods. While black-box interpretations assume no knowledge about the underlying model, white-box methods only work by using the model parameters.

This terminology of discriminating between bb and wb models may not be confused with the nature of the underlying models: Models still remain of black-box nature even though a white-box

method may contribute to making the decision making process of such a model more insightful.

The following section details the two categories and will give examples of the state-of-the-art interpretation methods within each group.

2.1 Black-box Methods

Black-box interpretations assume no knowledge about the model thus treating it as a black-box. The underlying model is approximated by learning its behavior with an interpretable model, e.g. a linear model. For this, the common approach is to approximate the relationship between the input samples and the corresponding prediction by the model. As the model itself does not need to be known for using such a model-agnostic approach, these approaches can be used in scenarios where the model itself is not directly accessible. A black-box interpretation offers the big advantage to be applicable to any model.

LIME. Local Interpretable Model-agnostic Explanations (LIME) [12] perturbs the input and observes how the predictions change. In image classification, LIME creates a set of perturbed instances by dividing the input image into interpretable components (contiguous superpixels), and runs each perturbed instance through the model to get a probability.

Figure 2 shows an example of LIME applied to an image and the neural network model Inception-v3 [16].




Figure 2: Visualization of LIME.

SHAP. SHAP, short for SHapley Additive exPlanations, calculates an additive feature importance score for each prediction with a set of desirable properties, such as local accuracy or consistency that its antecedents lacked.

2.2 White-box Methods

On the other side are white-box methods, where the underlying model is known with all its parameters. Thus, the interpretation can be directly computed by using the model instead of relying on an approximation of f_N as within the black-box methods. These methods typically rely on the relationship between an input sample, the underlying model's prediction and the associated activations of the model's hidden layers. Most methods in this area aim to highlight the features of the input that are important for the prediction. An example for model-transparent methods are gradient based and saliency map based methods.

Notations. A white-box interpretation method, in the following named interpreter \mathcal{I} , generates a heatmap

$$h_c^{\mathcal{I}}(\omega) = \mathcal{I}(x, c; \omega)$$

for a neural network with parameters ω and class c . The heatmap is a vector $h_c^{\mathcal{I}}(\omega) \in \mathbb{R}^d$, where the j -th value $h_{c,j}^{\mathcal{I}}(\omega)$ represents the importance score of the j -th input feature x_j for the final score of class c .

Layer-wise Relevance Propagation (LRP) relies on a Taylor series close to the prediction point rather than partial derivatives at the prediction point itself

Gradient-weighted Class Activation Mapping (Grad-CAM) [14] To further improve the quality of the visualization, attribution methods such as heatmaps, saliency maps or class activation methods (GradCAM[292]) are used. Grad-CAM uses the gradients of any target concept, flowing into the final convolutional layer to produce a coarse localization map, highlighting the important regions in the image for predicting the concept

SimpleGrad (SimpleG)

3 MANIPULATION METHODS

The focus of most works is on adversarial attacks against the prediction of a machine learning model. Contrary, the focus of this paper is on the attacks on the interpretation without changing the prediction.

A manipulation method refers to a method influencing an interpretation method \mathcal{I} to yield a wrong interpretation. This influence on the interpretation method is also called *fooling* or an *attack*.

The goal is to apply efficient and for humans imperceptible perturbations to change the interpretability of the input sample while leaving the prediction unchanged.

Formally, the problem can be defined as: **Definition 2: Interpretation Manipulation Method**

A manipulation method \mathcal{F} is defined as a method for altering the output of an explanation method while leaving the model performance of the neural network N roughly unchanged. This must be the case, as the attack is targeted to fool the explanation method, and is essentially not targeted to fool the model itself. Fooling the model would only disclose the vulnerability of the model but would not allow to gain insight into the stability of the fooling method. Hence, the main criteria for an explanation manipulation method must be fulfilled:

1. The model's prediction stays approximately the same, i.e. $f_N(x_{adv}) \approx f_N(x)$ or $f_{N_{adv}}(x) \approx f_N(x)$
2. The explanation map $h(x_{adv})$ is significantly different to the explanation map resulting from non adversarial models or inputs $h(x)$, i.e. $h(x_{adv})$ or

$$\arg \max_{\delta} \mathcal{D}(\mathcal{I}(x_i, \omega), \mathcal{I}(x_i + \delta, \omega))$$

3. In case the attack is in the input domain of the model, the perturbation of input samples must be imperceptible by humans. According to [5], the norm of the added perturbation δ to an input sample x thus must be small, i.e. $\|\delta x\| = \|x_{adv} - x\|_{\text{inf}} \leq \epsilon$

Other criteria to a manipulation method are the following:

Effectiveness. The manipulation scheme is inexpensive to conduct. Input manipulations are by definition inexpensive, as the perturbation can be applied to single input samples. Model manipulations are more expensive as they require the model parameters to be adapted. However, an adversarial model can be obtained by fine-tuning the model with an adapted objective function. This fine-tuning also has the advantage that the model is adapted to include a systematic bias and can thus be applied to fool explanation methods without further adapting the model or input samples.

Transferability. The manipulation does not only fool one type of explanation method, but its effect transfers to other types.

Most of the explanation methods outlined in sec. section 2 have been shown to be vulnerable to adversarial perturbations. Manipulation methods often show that there exist small feature changes resulting in a change of the explanation methods output while the output of the model itself does not change. Most approaches aim at providing a relevance measure of the input features.

3.1 Manipulation Levels

Adversarial Input Manipulation. The general approach is to perturb input data while observing the effect of this perturbation on the model prediction. As found in TODO, visually-imperceptible perturbations of an input image can make explanations worse for the same model and interpreter.

Adversarial Model Manipulation. Contrary to input manipulations, model manipulations do not operate on the input space but rather on the model parameter space itself. As first introduced by Heo et al. [8] in 2017, this line of research is comparably new. Adversarial model manipulations are obtained by fine-tuning the model on the same data but with an adapted objective function. [8] propose the adapted loss function of

$$\mathcal{L} = \mathcal{L}_{CE}(\mathcal{D}; \omega) + \lambda \cdot \mathcal{L}(\mathcal{D}; \omega; \omega_0)$$

where \mathcal{L}_{CE} would be the standard cross-entropy classification loss.

The authors find that perturbed model parameters can also make explanations worse for the same input images and interpreters.

3.1.1 Transferability of Manipulations. [8] find that fooling one explanation method with a fooling scheme transfers to other methods.

3.2 Manipulation Targets

A further categorization of interpretation methods can be made based on the target of the manipulation.

Untargeted Manipulations. The majority of manipulations is untargeted, meaning that the applied perturbations are mostly random and not designed to change the prediction for a specific portion of an input sample.

Targeted Manipulations. On the contrary, targeted manipulations are designed to specifically alter the explanation of a distinct portion. This specific portion might be an object in the input image in the context of image classification. For instance [8] introduce a fooling scheme in which the interpretations of the classes elephant and school bus are swapped. Manipulations on the level of the

model are mostly targeted, as the explanation methods are being fooled by adapting the model parameters.

3.3 Evaluation Criteria

As outlined in section section 3, there exists a plethora of interpretation methods differing in the assumption about the model character and also in style how interpretations are obtained. Thus, reliable evaluation methods are required allowing for a choice of an appropriate and robust explanation method. Evaluations of the quality of an explanation method can be separated into qualitative and quantitative evaluations.

Quantitative Evaluations. As the goal of interpreter manipulations is to fool an interpreter, thus altering the output of an interpreter, it is straightforward to compare interpretations before and after perturbation [7]. Common metrics for quantitatively measuring this similarity between interpretations are the following:

- **Spearman's rank order correlation.** As interpretation methods rank the features based on their importance, the rank correlation [15] is a natural measure for comparing interpretations.
- **Intersection of the top- k features.** For some tasks, only the top- k features are relevant, such that a comparison between these top- k features is insightful.
- **Fooling Success Rate (FSR).** Heo et al. [8] introduce the concept of the Fooling Success Rate (FSR) with the aim to create a measure for systematically measuring the robustness of an interpretation method to adversarial model manipulations. The FSR measures on how many samples from a dataset an interpretation method \mathcal{I} is fooled by a model manipulation. The higher the FSR, the more often an interpretation method is fooled.

Qualitative Evaluations. As interpretations are attributed to input features, the resulting relevance values l can be easily mapped to the input vector x . Looking at these evaluations for specific samples is informative albeit not usable to obtain a general statistic about the explanation methods quality.

4 TRANSFERABILITY OF PERTURBATIONS

sec:transferability)

input perturbations do not propagate to the whole validation set. On the contrary, model manipulations are non-local perturbations, meaning that they do not merely perturb an input sample but rather effect all samples in the way that the model itself is changed.

5 EXPLAINING MANIPULATIONS

There is an abundance of examples and scenarios where model explanations fail. However, there are few papers specifically targeting why these manipulations work in the first place.

6 UNCOVERING MANIPULATIONS

7 SANITY CHECKS FOR INTERPRETERS

Checking the robustness, scope and hence the quality of model interpreters has become an indispensable step in explainable machine learning.

[1] propose a number of randomization tests for saliency-map based interpreters. The authors find that most methods fail their tests, and warn of the danger of visual assessment.

Additionally, in order to account for adversarial model manipulations, Heo et al. [8] propose to expand the criteria for checking the robustness of interpreters further.

8 BENCHMARKING INTERPRETATIONS

Evaluating explanation and interpretation methods is difficult as ground truth is mostly lacking. In most applications, it is not known which input features are most important.

9 EXPERIMENTS

In this section, several experiments are evaluated that were conducted to replicate findings of other studies. Furthermore these approaches are extended to other domains and datasets.

9.1 Explanation Methods

9.2 Manipulation Methods

9.3 Models

9.4 Datasets

ImageNet German Loan Dataset

Recidivism Dataset

Adult Income Dataset The adult income dataset [6] contains .. samples.

10 DISCUSSION

This paper summarizes the current approaches to manipulating model interpretation methods. On the one hand, the findings suggest that our models are not fully aligned with how human information processing works. If machine learning models would decide by the criteria we humans employ for tasks such as image classification, there would be no fooling of interpretation models by input or model manipulations. On the other hand, it was shown that advances in machine learning models has led to models that rely too much on the data they are trained on, thus showing a high susceptibility to ood properties or properties that are highly correlated with labels in the dataset but are not distinctive in the real world (such as image backgrounds). Models and interpreters can still be misled in a large and systematic manner.

However, we hope that the vastly expanding and progressing field of XAI will help to move towards more robust, reliable and human-aligned machine learning models.

10.1 Conclusion

Finally, it must be noted that the suitability of a method depends on its application domain.

Much critique has been applied to methods aiming at interpreting complex and potentially non-interpretable models. Some researchers argue it is not worthwhile to study non interpretable systems while dismissing that using inherently interpretable models in the first place might be the better approach.

Adversarial attacks show that machine learning systems are still fundamentally fragile: They may be successful in a number of

tasks, but fail to adapt to ood scenarios, i.e. when being applied to unfamiliar territory.

The findings about manipulating interpretations do not suggest that interpretations are completely meaningless, just as adversarial attacks on predictions models do not imply that machine learning models are useless. However, they suggest that there still are fundamental flaws in the way neural networks operate and that much caution and supervision should be applied if they are to be deployed in the real world.

do not suggest that interpretations are meaningless, just as adversarial attacks on predictions do not imply that neural networks are useless.

This paper follows the footsteps of [9], trying to caution against blindly putting faith into post-hoc explanation methods. Moreover, we propose that checking the robustness of interpretation methods not only with respect to adversarial input manipulations but also with respect to adversarial model manipulation should be an proof of concept.

10.2 Future Work

We see several possible future directions of future work. Firstly, for approaching the discrepancy of in-lab and real-life applications of, more focus might be laid in the development better performance metrics for both measuring the performance of machine learning models as well as their interpreters. More specifically, it might be fruitful to further investigate the correlation between ood samples and the performance of an interpretation method.

There is also no work proposing a benchmarking for ...

REFERENCES

- [1] Julius Adebayo, Justin Gilmer, Michael Muelly, Ian Goodfellow, Moritz Hardt, and Been Kim. 2018. Sanity checks for saliency maps. *Advances in neural information processing systems* 31 (2018), 9505–9515.
- [2] Alejandro Barredo Arrieta, Natalia Díaz-Rodríguez, Javier Del Ser, Adrien Benetot, Siham Tabik, Alberto Barbado, Salvador García, Sergio Gil-López, Daniel Molina, Richard Benjamins, et al. 2020. Explainable Artificial Intelligence (XAI): Concepts, taxonomies, opportunities and challenges toward responsible AI. *Information Fusion* 58 (2020), 82–115.
- [3] Titus J Brinker, Achim Hekler, Alexander H Enk, Carola Berking, Sebastian Haferkamp, Axel Hauschild, Michael Weichenthal, Joachim Klode, Dirk Schadendorf, Tim Holland-Letz, et al. 2019. Deep neural networks are superior to dermatologists in melanoma image classification. *European Journal of Cancer* 119 (2019), 11–17.
- [4] Alexandra Chouldechova. 2017. Fair prediction with disparate impact: A study of bias in recidivism prediction instruments. *Big data* 5, 2 (2017), 153–163.
- [5] Ann-Kathrin Dombrowski, Maximilian Alber, Christopher Anders, Marcel Ackermann, Klaus-Robert Müller, and Pan Kessel. 2019. Explanations can be manipulated and geometry is to blame. In *Advances in Neural Information Processing Systems*. 13589–13600.
- [6] Dheeru Dua and Casey Graff. 2017. UCI Machine Learning Repository. <http://archive.ics.uci.edu/ml>
- [7] Amirata Ghorbani, Abubakar Abid, and James Zou. 2019. Interpretation of neural networks is fragile. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 33. 3681–3688.
- [8] Juyeon Heo, Sunghwan Joo, and Taesup Moon. 2019. Fooling Neural Network Interpretations via Adversarial Model Manipulation. *CoRR* abs/1902.02041 (2019). [arXiv:1902.02041](http://arxiv.org/abs/1902.02041) <http://arxiv.org/abs/1902.02041>
- [9] Zachary C Lipton. 2018. The mythos of model interpretability. *Queue* 16, 3 (2018), 31–57.
- [10] W James Murdoch, Chandan Singh, Karl Kumbier, Reza Abbasi-Asl, and Bin Yu. 2019. Definitions, methods, and applications in interpretable machine learning. *Proceedings of the National Academy of Sciences* 116, 44 (2019), 22071–22080.
- [11] Thai-Son Nguyen, Sebastian Stueker, and Alex Waibel. 2020. Super-Human Performance in Online Low-latency Recognition of Conversational Speech. *arXiv preprint arXiv:2010.03449* (2020).

- [12] Marco Tulio Ribeiro, Sameer Singh, and Carlos Guestrin. 2016. "Why should I trust you?" Explaining the predictions of any classifier. In *Proceedings of the 22nd ACM SIGKDD international conference on knowledge discovery and data mining*. 1135–1144.
- [13] Robin Ruede, Verena Heusser, Lukas Frank, Alina Roitberg, Monica Haurilet, and Rainer Stiefelhagen. 2020. Multi-Task Learning for Calorie Prediction on a Novel Large-Scale Recipe Dataset Enriched with Nutritional Information. *arXiv preprint arXiv:2011.01082* (2020).
- [14] Ramprasaath R Selvaraju, Michael Cogswell, Abhishek Das, Ramakrishna Vedantam, Devi Parikh, and Dhruv Batra. 2017. Grad-cam: Visual explanations from deep networks via gradient-based localization. In *Proceedings of the IEEE international conference on computer vision*. 618–626.
- [15] Charles Spearman. 1961. The proof and measurement of association between two things. (1961).
- [16] Christian Szegedy, Vincent Vanhoucke, Sergey Ioffe, Jon Shlens, and Zbigniew Wojna. 2016. Rethinking the inception architecture for computer vision. In *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2818–2826.
- [17] Alexander Warnecke, Daniel Arp, Christian Wressnegger, and Konrad Rieck. 2020. Evaluating explanation methods for deep learning in security. In *2020 IEEE European Symposium on Security and Privacy (EuroS&P)*. IEEE, 158–174.
- [18] Kai Xiao, Logan Engstrom, Andrew Ilyas, and Aleksander Madry. 2020. Noise or signal: The role of image backgrounds in object recognition. *arXiv preprint arXiv:2006.09994* (2020).